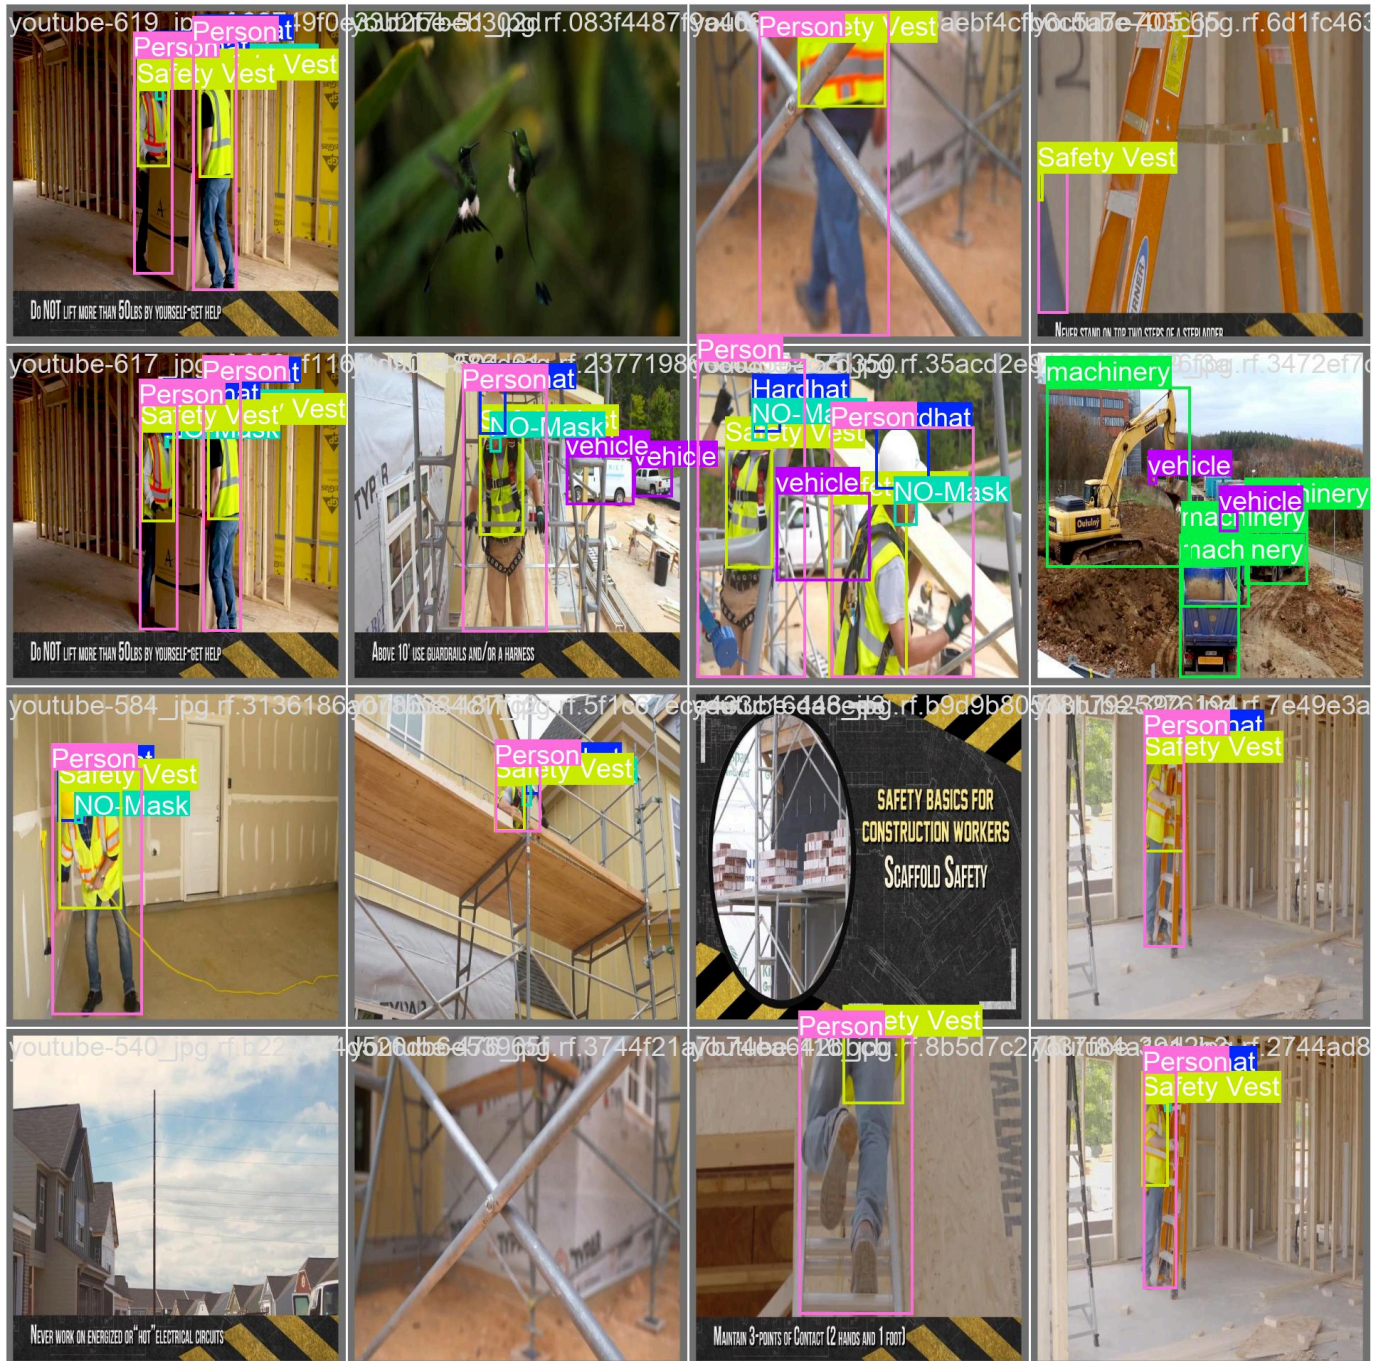


YOLO Model Fine-Tuning and Inference

Overview

This repository contains a Python-based framework for fine-tuning and performing inference with YOLO (You Only Look Once) models. The script provides functionalities such as dataset preparation, downloading model weights, displaying images, and computing dataset statistics.



Features

- Configurable training and inference settings
- Automatic dataset YAML file generation
- Image processing utilities (displaying and analyzing images)
- Dataset statistics computation
- Base model inference using pre-trained YOLO models
- Custom dataset support for training YOLO models
- Model evaluation and visualization tools

Dataset link: <https://universe.roboflow.com/roboflow-universe-projects/construction-site-safety>

Dependencies

Ensure you have the following Python libraries installed:

```
pip3 install -U -r requirements-dev.txt
```

Required Libraries:

- Python 3.x
- OpenCV (`cv2`)
- PyYAML (`yaml`)
- Pandas (`pandas`)
- Matplotlib (`matplotlib`)
- Seaborn (`seaborn`)
- Pillow (`PIL`)
- Ultralytics (`ultralytics`)
- Requests (`requests`)
- NumPy (`numpy`)

Configuration

The `CONFIG` class defines the core parameters:

- `DEBUG`: Enable debug mode (faster runs with fewer epochs)
- `DISPLAY_IMAGES`: Show images during processing
- `FRACTION`: Fraction of data used for debugging
- `CLASSES`: List of class names for detection
- `EPOCHS`: Number of training epochs
- `BATCH_SIZE`: Training batch size
- `BASE_MODEL`: YOLO model type (e.g., `yolov8s`)
- `CUSTOM_DATASET_DIR`: Path to dataset directory
- `OUTPUT_DIR`: Path for model outputs

- `IMAGE_SIZE`: Image resolution for model input

YOLO and Ultralytics

YOLO (You Only Look Once) is a cutting-edge object detection architecture known for its speed and accuracy. The idea is simple yet powerful—YOLO uses a single neural network to predict bounding boxes and class probabilities directly from full images in one evaluation, which makes it incredibly efficient for real-time detection.

Key Points:

- **Pretrained Weights:**

YOLO comes with various pretrained weights (like `yolov8n`, `yolov8s`, `yolov8m`, etc.) that have been trained on large, diverse datasets. These models provide a strong starting point, saving you time and computational resources compared to training from scratch.

- **Fine-Tuning:**

Instead of building a model from zero, you can fine-tune these pretrained weights on your own custom datasets. This is particularly useful for specialized tasks where the objects of interest may not be covered by the original training data.

- **Custom Annotations:**

Your dataset should include a `labels` folder containing annotations (bounding boxes around objects). These annotations are crucial for the model to learn how to detect and classify objects relevant to your specific application.

- **Ultralytics Library:**

Ultralytics offers a user-friendly interface to work with YOLO models. It streamlines the process of training, fine-tuning, and performing inference. When fine-tuning, you load the pretrained weights and adjust them using your custom data and labels.

- **YAML Configuration File:**

Fine-tuning with Ultralytics requires a YAML file that specifies your dataset configuration. This file should include:

- `names`: A list of class names.
- `nc`: The number of classes.
- `train`: The path to the training images.
- `val`: The path to the validation images.
- `test`: (Optional) The path to the test images.

For example:

```
# data.yaml
names:
  - Hardhat
  - Mask
```

```
- NO-Hardhat
- NO-Mask
- NO-Safety Vest
- Person
- Safety Cone
- Safety Vest
- machinery
- vehicle
nc: 10
test: data/test/images
train: data/train/images # Optional
val: data/valid/images
```

- **Best Practices:**

- **Documentation:** When generating code (especially with an LLM), always provide background context about what you're using and the expected output.
- **File Paths:** Avoid using absolute paths in Python projects. Instead, save files relative to the project's root directory for better portability and a cleaner structure.

The Anatomy of an o1 Prompt

I want a list of the best medium-length hikes within two hours of San Francisco.

Each hike should provide a cool and unique adventure, and be lesser known.

For each hike, return the name of the hike as I'd find it on AllTrails, then provide the starting address of the hike, the ending address of the hike, distance, drive time, hike duration, and what makes it a cool and unique adventure.

Return the top 3.

Be careful to make sure that the name of trail is correct, that it actually exists, and that the time is correct.

--

For context: my girlfriend and i hike a ton! we've done pretty much all of the local SF hikes, whether that's presidio or golden gate park. we definitely want to get out of town -- we did mount tam pretty recently, the whole thing from the beginning of the stairs to stinson - it was really long and we are definitely in the mood for something different this weekend! ocean views would still be nice. we love delicious food. one thing i loved about the mt tam hike is that it ends with a celebration (Arriving in town to breakfast!) The old missile silos and stuff near Discovery point is cool but I've just done that hike probably 20x at this point. We won't be seeing each other for a few weeks (she has to stay in LA for work) so the uniqueness here really counts.

Goal

Return Format

Warnings

Context Dump

Code Example: Fine-Tuning with YOLO and Ultralytics

Below is a sample code snippet showing how to fine-tune a YOLO model using the Ultralytics library:

```
from ultralytics import YOLO

# Load the pretrained YOLO model (e.g., yolov8s)
model = YOLO("models/yolov8s.pt")

# Fine-tune the model on your custom dataset
# Note: The 'data.yaml' file should include keys for names, nc, train, val, and optionally test.
#       Set task="detect" for object detection tasks.
results = model.train(data="data.yaml", task="detect", epochs=50, batch=16)

# After training, you can run inference on new images
results = model("path/to/your/image.jpg")
```

This example demonstrates how to load a pretrained model and fine-tune it using a YAML file that holds the dataset configuration. The `data` parameter in the `model.train()` call is set to `"data.yaml"`, and the `task` parameter is set to `"detect"`, ensuring the model is properly set up for object detection tasks.

Usage

1. Install Dependencies

```
pip install -U -r requirements-dev.txt
```

2. Run the Model

```
python3 run.py
```

3. Modify Configuration

To change core settings, update the following values in the configuration file:

```
# Set to True for development mode (reduced dataset, fewer epochs)
# Set to False for production mode (full dataset, full training)
DEBUG = True

# Set to True to view images and plots during processing
# Set to False for faster processing without visual output
DISPLAY_IMAGES = False

# Automatically adjusts dataset size based on DEBUG mode
FRACTION = 0.05 if DEBUG else 1.0 # Uses 5% of data when DEBUG is True
```

```
# Automatically adjusts training epochs based on DEBUG mode
EPOCHS = 3 if DEBUG else 50 # Uses 3 epochs when DEBUG is True
```

4. View MLflow UI

After training is completed, run the following command to launch the MLflow UI:

```
mlflow ui --backend-store-uri runs/mlflow
```

Directory Structure

```
.
├── data/
│   ├── train/
│   ├── valid/
│   └── test/
├── models/
│   └── yolov8s.pt
├── output/
│   ├── runs/
│   │   └── mlflow/
│   └── logs/
└── script.py
```